


—●— R-REBEL ( $\ell_1$ -std)    -■- R-REBEL (Huber-std)    —▲— GRPO + entropy    —◆— GRPO, base-LM reference

(a) where each floor's point sits



(b) what the reward scores:  $P(c \geq \tau)$

